

UNWITTING WRONGDOERS AND THE ROLE OF MORAL DISAGREEMENT IN BLAME¹

(in David Shoemaker [ed], *Oxford Studies in Agency and Responsibility, Volume 1* [Oxford University Press, 2013], pp. 225–245)

Matthew Talbert
West Virginia University

In “Culpable Ignorance,” Holly Smith says that “[i]gnorance of the nature of one’s act is the pre-eminent example of an excuse that forestalls blame” (H. Smith 543). So, to adapt an example of Smith’s, a doctor who gives a patient the wrong treatment might avoid blame if she thought she was providing her patient with the best care she could. Whether the doctor avoids blame in this way will depend, in part, on whether she was blameworthy for being mistaken about what treatment was best. If the doctor prescribed the wrong treatment because of an earlier “benighting act” that impaired (or failed to improve) her cognitive position, then, *if we take the doctor to be culpable for the benighting act*, we may also regard her as blameworthy for her ignorance and for the consequences of her ignorance (H. Smith 547).

Smith’s focus is on agents whose ignorance is circumstantial rather than moral: the doctor in the above example is unaware that the treatment she prescribes is not best for her patient, but she presumably is aware that she morally ought to give her patient the best treatment.² More recently, Gideon Rosen has applied an approach like Smith’s to cases of both

¹ I would like to thank the participants at the 2011 New Orleans Workshop on Agency and Responsibility for their questions and comments on this paper, particularly Sarah Buss, T. M. Scanlon, Sarah Stroud, and Gary Watson. I am also very grateful to David Shoemaker and to two anonymous referees for Oxford University Press for their written comments.

² It is true that the doctor suffers from normative ignorance since she does not know that she has prescribed an inappropriate medication, but this is not what I mean by “moral ignorance.” A doctor who suffers from moral ignorance (as I shall use the phrase) would not know that it is wrong to intentionally prescribe a patient the wrong medication. The doctor in the example suffers from circumstantial ignorance because she is not aware that, given her circumstances, her action will lead to bad consequences for her patient.

circumstantial and moral ignorance (Rosen 2004).³ Rosen assumes that for an agent to be blameworthy for actions that issue from moral or circumstantial ignorance, the agent must be culpable for her ignorance. This assumption plays a crucial role for Rosen in a skeptical argument that calls into question many intuitive judgments of moral blameworthiness.

Though I reject Rosen's conclusion, I find his skeptical argument powerful. Indeed, and as I attempt to show, the argument is sufficiently strong that a relatively conservative approach to overturning it does not succeed.⁴ Instead, if we are to avoid the skeptical conclusion, we must reject the plausible sounding assumption that unwitting wrongdoers are blameworthy only if they are culpable for their ignorance. To this end, I argue that while ignorance of the circumstances and consequences of one's actions often undermines blame, moral ignorance typically does not do so. For example, while I might not be blameworthy for injuring you if I was unaware that my action would have that result, I likely would be blameworthy if I were simply unaware that injuring you is impermissible. I will argue, moreover, that a morally ignorant wrongdoer can be blameworthy even if it is not her fault that she is ignorant of the moral status of her behavior, and even if it would be unreasonable to expect her to be aware of its status. In the context of making these points, I also try to shed light on the role that moral disagreement plays in our judgments of blameworthiness.

1. Skepticism about Moral Responsibility

Gideon Rosen has advanced a skeptical perspective on moral responsibility based on the assumption that a wrongdoer is excused if she is non-culpably unaware that she does wrong. For Rosen, moral responsibility (for bad actions) amounts to being liable to the "sanctions"

³ Also see Rosen (2003). Michael Zimmerman (1997) offers a view similar to Rosen's; also see Zimmerman (2008).

⁴ The conservative approach I have in mind is the one pursued by William FitzPatrick (2008). I discuss FitzPatrick's view in Section 2.

associated with moral blame; so, for Rosen's purposes, "we may simply identify moral responsibility with culpability or blameworthiness" (Rosen 2004 296). As Rosen puts it, "[s]kepticism about moral responsibility is thus the thesis that *confident positive judgments of blameworthiness are never justified*" (Rosen 2004 296).

Like many philosophers working on the subject, Rosen takes moral blame to involve certain emotional responses. For example, "you blame X for doing A, when you resent him or feel indignant towards him for having done it" (Rosen 2004 296). Thus, if we judge that X is blameworthy, this means "that X is *liable* [in our judgment] to a negative emotional response of this sort for having done A, or equivalently, that some such response would be *appropriate* or *fitting*" (Rosen 2004 297). Rosen's skepticism amounts, then, to skepticism about the aptness of negative reactive attitudes like resentment. In what follows, I will understand moral responsibility in just the way Rosen suggests: the primary question will be about whether judgments of blameworthiness, and responses like resentment, are apt.

Rosen's skeptical argument proceeds as follows. Wrongdoing is either witting or unwitting. While a *knowing* wrongdoer may be directly blameworthy for her behavior, an unwitting wrongdoer will be blameworthy for her wrongdoing only derivatively or indirectly. This means that an ignorant wrongdoer will be morally responsible, in the sense of being open to blame, only if she is culpable for her ignorance.⁵ However, according to Rosen, "[i]gnorance is culpable only if it derives from culpable recklessness or negligence in the management of one's opinion" and this prior recklessness will be culpable only if it was knowing or was itself the product of knowing mismanagement of one's opinions (Rosen 2004 302). So, as Rosen puts it, ignorance is culpable—and unwitting wrongdoing blameworthy—only if it results, at some point, from an akratic act in which the agent knowingly violated an epistemic or moral duty:

⁵ Following the authors I discuss, I use "culpable" and "blameworthy" interchangeably.

One is responsible for the act done from ignorance only if one is independently responsible for something else. . . . this entails that *the only possible locus of original responsibility is an akratic act*. . . . Our first sin must be a knowing sin—a sin done in full knowledge of every pertinent fact or principle (Rosen 2004 307).⁶

Rosen applies this perspective to cases of circumstantial ignorance (like that of the doctor in the introduction), but also to cases of moral ignorance. For example, Rosen's view suggests that an "ancient slaveholder who . . . believes that it is morally permissible for him to buy and sell" slaves is blameworthy for his behavior "only if he is culpable for the moral ignorance from which he acts" (Rosen 2004 304).

The skeptical force of this position emerges when we consider the possibility that many unwitting wrongdoers have never committed the kind of akratic act that would make them culpable for their ignorance. Imagine an "ambitious capitalist who is mistaken about where to draw the line between permissibly aggressive business practices and reprehensibly ruthless business practices" (Rosen 2004 305). Rosen supposes (and I agree) that it is not hard to imagine a case in which the capitalist does something wrong without recognizing that his behavior is wrong and that he has been as reflective and careful as it is reasonable to expect him to be. Perhaps the capitalist's moral education was deficient, or perhaps the case he considers is "just a hard case and after thinking about it for a decent interval he has simply arrived at the wrong answer" (Rosen 2004 305). If we agree that "in reaching his [moral] conclusion our capitalist has not been reckless or negligent," Rosen thinks we should also agree "that his moral ignorance is not his fault" and that it would therefore "be a mistake to blame him for the wrong he does" (Rosen 2004 305).

⁶ Zimmerman's view is similar: if "one is culpable for ignorant behavior, then one is culpable for the ignorance to which this behavior may be traced" (Zimmerman 1997 418). However, "one is never in direct control of whether one is ignorant," so culpability for ignorant behavior must be indirect, which "presupposes direct culpability for something else. . . . Hence all culpability can be traced to culpability that involves lack of ignorance" (Zimmerman 1997 418).

Another example features Bill, who decides to tell a self-serving lie to his wife even though he “knows that it’s just plain wrong to lie to your wife” (Rosen 2004 305). Again, it is possible that Bill had a defective moral education that led him to believe, “through no fault of his own, that while moral considerations have some weight, they are not in general decisive” (Rosen 2004 305).⁷ Bill may thus be “blamelessly (though mistakenly) convinced that the balance of reasons comes down in favor of lying” (Rosen 2004 306). Rosen argues that in this case it would be unreasonable to blame Bill for telling the lie. Rosen asks: “Does it make sense to subject someone who blamelessly believes that he should do A, and then does it, to moral sanctions—to recrimination, resentment, righteous anger, contempt?” (Rosen 2004 306). Rosen thinks this doesn’t make sense. Instead, Bill’s wife should respond this way:

Poor Bill. Through no fault of his own he found himself believing that all things considered, he should lie. Given that he found himself in that state, I can hardly fault him for lying. Holding the judgment fixed, the lie itself was a perfect manifestation of practical rationality. I can fault him for the lie only if I can fault him for believing that in the circumstances, his selfish interests were more important than my moral interests. Since by hypothesis, it is not his fault that he held this view, I have no option but to conclude that he is not properly culpable for his bad action (Rosen 2004 306).⁸

Rosen’s official skeptical conclusion is that “*it would be unreasonable to repose much confidence in any particular positive judgment of responsibility*” (Rosen 2004 308). This follows from the claim that an unwitting wrongdoer is blameworthy only if she has an akrasia-involving, “*inculcating history*” (Rosen 2004 309) together with the claim that “it is almost always unreasonable to place significant confidence in” (Rosen 2004 308) the judgment that an agent has such a history. In a recent response to Rosen, William FitzPatrick has argued that in fact we often have good grounds for attributing akrasia to ourselves and others (FitzPatrick 593-99).

⁷ Unlike the ambitious capitalist, Bill is mistaken about the force of moral considerations rather than about a moral principle.

⁸ For discussion of a related example, see Rosen (2008 605-609).

However, as FitzPatrick notes, this leaves untouched Rosen's conclusion that blameworthy wrongdoing is always either knowing or the result of knowing wrongdoing (FitzPatrick 599). We are left, then, with a striking skeptical challenge, for very many ordinary wrongdoers may not have a relevant instance of akrasia in their past and so are not morally responsible for their wrongdoing.⁹ Below, I consider Fitzpatrick's attempt to meet this skeptical challenge.

2. FitzPatrick's Reply to Rosen

FitzPatrick accepts the "intuition at the core of [Rosen's] argument, that it is unfair to blame someone for an action done out of ignorance that he cannot fairly be blamed for having" (FitzPatrick 601). However, FitzPatrick rejects Rosen's claim that unwitting wrongdoers are culpable for their ignorance only if it resulted from knowing wrongdoing. According to FitzPatrick, an unwitting wrongdoer will also be culpable for her ignorance if she "could reasonably have been expected to take measures that would have corrected or avoided it" (FitzPatrick 609).

To make his case, FitzPatrick expands on Rosen's "ambitious capitalist" example and considers Mr. Potter, the ruthless businessman in Frank Capra's 1964 film, *It's a Wonderful Life*. Potter is callous, vindictive, greedy, and dishonest, but let us suppose that he wrongly believes that his behavior is permissible. FitzPatrick argues that even if Potter did not akratically acquire his false moral beliefs, there are three central factors "that make most of us confident that Potter's moral ignorance is culpable" (FitzPatrick 605). First of all, Potter never engaged in the kind of moral reflection that might have helped him correct his moral ignorance, yet "[t]here were no relevant limitations in his social context or in his capabilities" that would have made such reflection unreasonably difficult (FitzPatrick 605). This suggests that "[t]he failure of

⁹ This is close to the skeptical conclusion advanced by Zimmerman (1997 425).

adequate reflection” on Potter’s part “was instead the result of voluntary exercises of vices such as overconfidence, arrogance, dismissiveness, laziness, dogmatism, incuriosity, self-indulgence, contempt, and so on” (FitzPatrick 605). According to FitzPatrick, Potter “could thus reasonably have been expected to take steps that would have eliminated that [moral] ignorance, by refraining from exercising those vices and instead taking advantage of” available opportunities for moral and epistemic improvement (FitzPatrick 605).

So on FitzPatrick’s view, given Potter’s general capacities, and the fact that “the opportunity for improved normative understanding was clearly present in his social context,” it is reasonable to hold Potter to the expectation that he correct his moral ignorance (FitzPatrick 603-4). Thus, Potter is culpable for his ignorance and blameworthy for behavior that flows from it.

I agree with FitzPatrick that Potter is blameworthy for his bad behavior, but I am not convinced that it is reasonable to expect Potter to correct his moral ignorance. As FitzPatrick sees it, explaining Potter’s moral ignorance in terms of his vices helps us see that this moral expectation is reasonable: Potter’s failure to engage in reflection wasn’t forced on him by circumstances or inability, it resulted from vicious choices that he was free to omit. I content, however, that this emphasis on Potter’s vices is unhelpful for the point FitzPatrick wants to make.

FitzPatrick characterizes Potter’s failure to correct his moral ignorance as “the result of voluntary exercises of vices” (FitzPatrick 605). But what does Potter do *voluntarily*? He does not voluntarily behave in ways that he regards as impermissibly vicious—FitzPatrick’s point hangs on Potter not being a *knowing* wrongdoer. Rather, Potter voluntarily exercises his vices in the sense that he acts voluntarily and his vices shape his practical judgments in characteristic ways. Now if Potter’s vicious actions are voluntary in this sense, then he might have omitted them if he

judged himself to have reason to do so, but one symptom of Potter having his vices is that he sees little in favor of such an omission. If Potter is afflicted with all the vices FitzPatrick mentions, he is not likely to see much in favor of exploring the opportunities for moral improvement available to him, so perhaps it is unreasonable to expect him to do so.

Neil Levy argues similarly against FitzPatrick that it is not reasonable to expect Potter to correct his vices because it is likely not subjectively rational for Potter to do so. When we think about what can reasonably be expected of a person, Levy says we should consider what that person can do by way of a rational reasoning procedure, and what an agent can do in this way is a function of her internal reasons, her “actual representations and proattitudes” (Levy 2009 736). As Levy notes, even if Potter’s failure to subject his values to scrutiny is a manifestation of epistemic and moral vices,

by his lights, Potter governs his normative views adequately. He gives competing views the attention he takes them to deserve. . . . But if Potter does not see that he is managing his moral views badly, he has no (internal) reason to manage them any differently. Potter exhibits epistemic vices aplenty, but because he does not conceive of them as vices, he has no internal reason to refrain from so doing (Levy 2009 737).

Levy concludes that since Potter “could not rationally have taken advantage of the opportunities for moral improvement” with which he was presented, “we cannot reasonably expect him to do so” (Levy 2009 735).¹⁰

¹⁰ This discussion is taken up again in (Levy 2011 124-128). A referee for Oxford University Press points out that we use the word “expect” in normative and descriptive senses (to use the referee’s terminology). The descriptive sense presumably has to do with what we anticipate from an agent in view of his capacities and the context in which he acts, whereas the normative sense has to do with standards against which we measure agents. When I say, “I expect it to rain tomorrow,” I use “expect” in the descriptive sense. I anticipate rain; it seems like a safe bet. Neither Levy nor FitzPatrick expect moral reform from Potter in this sense; neither regards it as a safe bet. (Though FitzPatrick is keen to draw attention to Potter’s abilities and his social context and the way that these make reform at least possible; he regards such reform as a live possibility even if not as a likely outcome.) By contrast, both Levy and FitzPatrick agree that Potter falls short of our expectations in at least one normative sense: they both regard Potter’s behavior as morally subpar. The question is whether this subpar behavior opens Potter up to moral blame. Here we might identify a slightly different sense in which “expect” can be normative, the sense in which we talk about *holding* someone to an expectation. Fitzpatrick thinks it is reasonable to hold Potter to our moral expectations because it is possible for him to live up to these expectations given his abilities and the context in which he acts.

There is also the question of how Potter acquired his vices. As FitzPatrick notes, it may seem “problematic that the vices Potter exhibits in his epistemically debilitating choices may trace back to his childhood and may be largely a result of moral (bad) luck” (FitzPatrick 607). FitzPatrick deflects this worry by noting that for most people “character traits are not merely given but are formed, reformed and continuously shaped by our choices from the point of moral maturity onward” (FitzPatrick 608). But this is not a helpful response to the problem of constitutive moral luck since it simply pushes the problem back to earlier stages of Potter’s development. Apparently, Potter made poor choices as he shaped his character. But why did he make such poor choices? By hypothesis, he did not *knowingly* make poor character-forming choices, so perhaps Potter’s tendency to make “epistemically debilitating” choices is explained by a tendency to see poor self-forming choices as choice-worthy. But in this case, Potter’s self-forming choices would seem to be shaped by incipient versions of the vices that his self-formation is invoked to explain. And given the presence of these incipient vices, why should we expect Potter to make the right choice when it comes to choosing whether to act in a way that will strengthen his vices?

Finally, it is worth noting that FitzPatrick says that “cultural and historical contexts” may make it unreasonable to expect a wrongdoer to know better because “the relevant [moral] knowledge isn’t reasonably available” (FitzPatrick 612). Aristotle, for example, may not be blameworthy for thinking slavery permissible because slavery would have presented “a genuinely hard case for someone in Aristotle’s circumstances” (FitzPatrick 600 note 24). Since social and cultural factors can turn a question that is easy *for us*—like “Is slavery wrong?”—into a difficult question for someone in ancient Greece, this undermines the culpability of the morally

Levy thinks it is unfair to hold Potter to our moral expectations because he can fulfill these only by behaving irrationally.

ignorant Greek. But Potter's vices seem to turn apparently easy moral questions into difficult ones, so perhaps these internal obstacles to moral knowledge should (at least on FitzPatrick's account) have the same excusing force as cultural impediments to moral knowledge.

3. Rejecting the Skeptical Argument

I agree with Levy that it is unreasonable to expect Potter to recognize and correct his moral ignorance. If this is right, and if it is also true that Potter never committed a relevant act of knowing wrongdoing, then he would not be culpable for his moral ignorance on either FitzPatrick's or Rosen's account of culpable ignorance. What goes for Potter presumably goes for many other unwitting wrongdoers, so if we accept the assumption—endorsed by Rosen, FitzPatrick, and Levy—that unwitting wrongdoers are blameworthy only if they are culpable for their ignorance, then we should conclude that many (perhaps very many) ordinary wrongdoers are not blameworthy for their bad behavior. Fortunately, we need not accept this conclusion because there is good reason to reject the assumption that unwitting wrongdoers are blameworthy only if they are culpable for their ignorance.

The central instance of wrongdoing in *It's a Wonderful Life* occurs when Mr. Potter keeps \$8000 that George Bailey's uncle Billy misplaced. Potter keeps the money, believing (and hoping) that this will cause the foreclosure of Bailey Building and Loan, as well as the prosecution of George Bailey for bank fraud, and that this will leave the town of Bedford Falls ripe for economic exploitation. Recall that we are concerned here with whether someone like Potter is morally responsible for his behavior in the sense of being an appropriate target for blaming attitudes like resentment. Plausibly, if George Bailey were to resent Potter, this resentment would be provoked by the fact that Potter willingly acted so as to cause a bad

outcome for George, and that he did so because of a desire to injure George and to extract an economic benefit from George's misfortune. These aspects of Potter's behavior make George's resentment natural and appropriate. If Potter's actions had lacked these features—if he had been coerced or had thought he was acting in George's best interests—then it would be inappropriate for George to resent Potter. But if Potter deliberately injured George, and did so for self-serving reasons, then George can quite reasonably point to these facts to explain and justify his resentment even if Potter happens to regard his behavior as permissible.¹¹

The features of Potter's behavior just mentioned make George's blaming responses appropriate because of the way these features are tied up with Potter's expression of contemptuous judgments and attitudes towards George.¹² For example, if Potter deliberately injured George for self-serving reasons, then this behavior expresses the implicit judgment that George's welfare is unimportant in comparison with whether Potter achieves his aims. (In fact, it's not just that George's welfare is unimportant for Potter, he regards the possibility of injuring George as a reason for acting.) Given the contemptuous judgments that inform Potter's behavior, George can reasonably regard Potter's behavior as unjustifiable and morally offensive—rather than as merely harmful or unwelcome—and thus as proper grounds for resentment.

Similar points apply to Rosen's example of Bill. Recall that, through no fault of his own, Bill believes that his selfish reasons for lying to his wife trump the moral reasons against doing so. Rosen says that

[a]nyone who bears our principles [about culpable ignorance] in mind and nonetheless judges that Bill is responsible for lying to his wife, is thereby committed to the view that somewhere in the story of that lie there exists a full-blown episode of altogether knowledgeable wrongdoing (Rosen 2004 307-8).

¹¹ In Section 5, I address the concern that Potter's non-culpable moral ignorance makes it unfair to blame him.

¹² For other applications of this kind of approach to blameworthiness, see T. M. Scanlon (1998) and (2008), Angela Smith (2005) and (2008), and my own (2008) and (2012a).

“Responsibility” is an elastic term, so we can interpret the claim, “Bill is responsible for lying to his wife,” in different ways. By hypothesis, Bill is not at fault for thinking that he has decisive reason to lie, so Bill is not responsible for lying to his wife in the sense of having played a particular sort of causal role in the process that led him to think that lying is the thing to do. But whether Bill is responsible for lying in this sense is not what is at issue. Again, the issue is whether Bill is morally responsible in the sense that attitudes like resentment would be a fitting response to him. The central question about Bill is whether it makes “sense to subject someone [like Bill] who blamelessly believes that he should do A, and then does it, to moral sanctions—to recrimination, resentment, righteous anger, contempt?” (Rosen 2004 306).

I take it that Rosen believes that Bill is responsible for lying, in the sense of being open to attitudes like resentment, only if he is responsible for lying in the sense of having played the right sort of role in bringing about his own tendency to favor lying. But Bill’s case is a good example of why moral responsibility, in the sense of blameworthiness, does not require that agents have this sort “inculcating history.” It is reasonable for Bill’s wife to blame him because of the way his lying expresses Bill’s morally faulty judgments and attitudes. Bill has these faults, and they contribute to his behavior, regardless of whether he is *at fault* for having them, and regardless of whether Bill genuinely believes that lying to his wife is the thing to do.

One of Bill’s faults has to do with how his wife’s interests rate when he is trying to figure out what to do: Bill is willing to overlook his wife’s interests when they conflict with Bill avoiding trouble. It is reasonable, then, to attribute to Bill the judgment that his wife’s interest in not being lied to can be overlooked, if that is how Bill can get what he wants. If Bill’s wife were to find out how her interests rate with her husband, that he lied to her and the basis on which he did so, then she would have good grounds for blame. That is, it would be appropriate for her to

be offended and hurt by what Bill's action expresses, for her to protest that her interests ought to rate more highly with Bill, to resent him for his callousness, to insist that he change his ways, and so on.

The general perspective developed above is as follows. Even if a wrongdoer is ignorant of the fact that her behavior is wrong, and even if this ignorance is not her fault, her actions may still express the contemptuous judgment that certain others do not merit consideration, that their interests do not matter, and that their objections can be overlooked. If one is injured by a wrongdoer who is moved by such judgments, then the attitudes and responses involved in moral blame are reasonable regardless of what the wrongdoer thinks about the moral status of her behavior. I will develop this perspective below, but there is enough here to see one way of rejecting the skeptical perspective on moral responsibility with which this paper began. Rosen's skepticism depends on assuming that unwitting wrongdoers are open to blame only if they are culpable for their ignorance. However, as I have argued, certain features of an unwitting wrongdoer's behavior can qualify her for blame regardless of whether she is culpable for her ignorance.

4. Moral Ignorance and Moral Disagreement

Angela Smith has applied a view like the one I just outlined to show that agents are often morally responsible for things that are not under their direct control: their desires and emotions, their advertences and inadvertences, and so on. We are responsible for these things, on Smith's view, because of the evaluative judgments they express and the importance of these judgments for our interpersonal relations.

Smith gives special attention to the fact that “we often take what a person notices and neglects to have an enormous amount of expressive significance” (A. Smith 2005 242). What a person notices attracts our attention because we assume a connection between what one notices and what one values. According to Smith, “if one judges some thing or person to be important or significant,” this should “have an influence on one’s tendency to notice factors which pertain to the existence, welfare, or flourishing of that thing or person” (A. Smith 2005 244). And if one fails to notice such factors, this “is at least some indication that she does not accept this evaluative judgment [about the thing’s importance]” (A. Smith 2005 244). As Smith says, I may fail to “notice when my music is too loud,” or that “my advice is unwelcome,” or that “my assistance might be helpful to others,” and *even if these failures are involuntary*, they may still indicate “that I do not judge your needs and interests to be important, or at least that I do not take them very seriously” (A. Smith 2005 244).

Though she does not describe them this way, the examples Smith cites are instances of circumstantial ignorance. On Smith’s view, then, judgments about a circumstantially ignorant wrongdoer’s moral responsibility should track the plausibility of associating her behavior with interpersonally significant evaluative judgments—particularly judgments about the normative status of the needs and interests of those affected by the wrongdoer’s actions and omissions.

As should be clear from the previous section, I am largely in agreement with Smith, but it is worth emphasizing that we cannot always infer that an agent does not care about something from the fact that she fails to notice how her actions (or omissions) will affect it. Smith is aware of this; she notes that in some cases of inadvertence, “the person in question may be extremely tired or under a lot of stress” and this may “block the normal inference from what a person notices to what she cares about” (A. Smith 2005 244 note 14). Indeed, we sometimes hesitate to

make the inference from what a person notices to what she cares about even when we cannot point to stressors that intuitively explain an agent's inadvertence. Sometimes, and for no obvious reason, we just forget things, or fail to notice them, yet we may have as much concern for these things as we ought to have. This can be true even in cases in which an inadvertence has horrible consequences, such as when a parent mistakenly leaves a child in a hot car. In some of these cases, it is no doubt correct to infer that the parent has a condemnable lack of concern for the child's welfare. But in other cases, it is difficult to read the testimony of the parents involved and come away with the thought that their forgetting is best explained by a morally deficient degree of parental concern. Many of these parents seem to have been as concerned with their child's welfare as morality requires; yet, they left them behind all the same.¹³

We may often be unsure how to assess the sorts of cases just described, but my general point is uncontroversial: we should be cautious before we conclude that an unwitting wrongdoer's actions or omissions express evaluative judgments that might ground blame. More controversially, I would argue that such caution is particularly appropriate in cases in which the unwitting wrongdoer is ignorant of features of the context in which she acts, or of the likely consequences of her behavior. There is correspondingly less reason to be cautious about attributing blame-grounding judgments in cases of moral ignorance in which an agent is aware of the consequences of his behavior, but is unaware that it is wrong to bring about those consequences.

One consideration in favor of this last claim is obvious: if an agent is aware that act *A* will have consequence *C*, then, when she *A*'s, it is usually reasonable to attribute to her at least an implicit judgment about how the prospect of *C* bears on the question of whether to *A*. Of

¹³ Cf. Levy (2011 182). Gene Weingarten's (2008) Pulitzer Prize winning article, "Fatal Distraction," offers a detailed and affecting account of some of these cases.

course, the fact that an agent knows that her action will cause *C* does not entail that she formed an *objectionable* judgment about the significance of *C*. The reason we should expect the actions of morally ignorant wrongdoers like Mr. Potter to express objectionable judgments is that their moral ignorance is—unlike a lot of circumstantial ignorance—often a manifestation of a normative disagreement the agent has with those who object to his behavior. More specifically, the moral ignorance of such an agent is often tied up with a perspective that regards as unobjectionable the very thing that his victims take to make his behavior objectionable.

Mr. Potter's moral ignorance, for example, is partly constituted by the fact that he thinks he assigns George's interests their proper weight in his practical deliberations. Of course, we disagree with him about the normative significance of George's interests, and since we take our position to be the right one, we regard Potter as lacking moral knowledge. In such a case—where a wrongdoer's moral ignorance is part and parcel of a profound normative disagreement between us and him—it is no surprise that his behavior expresses judgments and attitudes that are, by our lights, objectionable.

Of course, our normative judgments may conflict more or less profoundly with the judgments we attribute to morally ignorant wrongdoers, and as the severity of this disagreement decreases, so too may the intensity of our blame. Take the case of Robert E. Lee. Those interested in burnishing Lee's reputation often note that he joined the secessionist cause in the U.S. Civil War out of loyalty to the state of Virginia. Lee did wrong, I assume, in leading Confederate troops against the Union, but he thought he was doing what duty and loyalty required. Suppose, however, that Lee was motivated to support the Confederacy solely by racial hatred and a desire to see slavery preserved. In this case, Lee would still be a morally ignorant wrongdoer, but his defenders would face a much more difficult task in convincing us to excuse

him. This is because Lee's actions would have been associated with judgments and attitudes that we find deeply objectionable. An excuse like, "Lee thought he was doing the right thing," would, I submit, have little influence on us if we found the judgments that informed his choices thoroughly repugnant. And if we do have some tendency to accept this excuse, I suspect it is because the judgments about the value and requirements of loyalty that purportedly guided Lee are not utterly foreign or repellant to us.

5. Internal Reasons and the Fairness of Blame

In this section, I respond to the claim that it is unfair to blame an unwitting wrongdoer who can't reasonably be expected to omit her bad behavior. Along the way, I develop the suggestion I made at the end of the last section that moral blame often rests on the recognition of a moral disagreement between ourselves and the one we blame. I will argue that what matters most for judgments of blameworthiness are the considerations that count as reasons for us (as issuers of blame), and not whether these considerations could have counted as reasons for those we blame.

Let us return to Neil Levy's discussion of Mr. Potter's blameworthiness. According to Levy,

[i]t is not reasonable to blame agents for actions they cannot (intentionally) omit by way of some reasoning procedure; we cannot hold them responsible for failing to do things they could do only by chance or through a glitch in their agency (Levy 2009 739).

For Levy, part of the problem with blaming Potter is that his actual proattitudes do not rationalize a decision to reassess his values and epistemic practices, and this means that he lacks a fair opportunity to avoid the actions for which we would blame him. As Levy sees it, "agents have a fair opportunity to avoid performing actions . . . only if they could have rationally chosen

to omit the action, and what agents can do rationally is a function of their internalist reasons alone” (Levy 2009 739).

I believe that the sense in which Potter has trouble avoiding wrong actions does not make it unfair to blame him. Let us agree that Potter has no internal reason to avoid a morally bad action, and that if he had avoided the action, this would have been because of a “glitch” in his agency, as Levy puts it. Importantly, this does not mean that Potter’s *actual* bad behavior is the result of a glitch, mistake, compulsion, or anything similar. If someone is subject to a compulsion or to “glitchy” agency, then some of her actions may be unavoidable and she will have a claim to being excused for this behavior. But the basis of excuse here—the reason the compulsive agent is not appropriately targeted with resentment—is that her behavior is not under her control in such a way that it can express the objectionable judgments that would make blame appropriate. Usually, an agent’s inability to avoid an action goes together with that action not being under the agent’s control in this way, but cases of moral ignorance like Potter’s are an exception. Even if Potter does not have rational access to doing the right thing, his bad behavior may be a manifestation of perfect reflective self-control and subjective practical rationality: he may be acting just as he likes and for reasons that really do speak in favor of so acting, given his proattitudes. But if Potter’s bad actions are guided in this way by his judgment about reasons, then they can express the kinds of attitudes and evaluative judgments that I have argued make blaming responses appropriate.¹⁴

In addition, Potter presumably has the capacity to avoid wrongdoing in the sense that, for any wrong action, he would have refrained from performing that action if he had taken himself to have a decisive reason to do so. Of course, if Potter’s actual constitution entails that he will see no such reason, then he will not rationally avoid his wrongdoing, and it is this lack of rational

¹⁴ There is an important relation here with the theme of Harry Frankfurt (1969). I develop this point in (2012a).

access to avoiding wrongdoing that Levy thinks makes blame unfair. But if Potter's bad actions are guided by his judgments about how to behave, and if he would have acted differently if he had recognized a reason to do so, then his wrongdoing is unavoidable mainly in the sense that an action that is wrong (by our lights) is bound to seem choice-worthy to Potter. Since this sort of unavoidability does nothing to make Potter's behavior less knowing, deliberate, and dismissive of the interests of those his actions affect, I do not see why it should make blame unfair.

So, while I agree with Levy that it is not reasonable to demand that Potter behave contrary to his internal reasons, what I think this tells us is that having subjective rational access to avoiding an action is not a requirement on being properly blamed for that action. In other words, an agent may be blameworthy for performing an action that she had no internal reason to avoid. In my view, approaches like Levy's focus too much on the perspective of the blamed agent, on what can count as a reason for him, and what can be expected of him. In fact, I think that our judgments about blameworthiness are often formed—and rightly so—on the basis of the normative considerations that *we* recognize, and not on whether those we blame could have assigned the same normative weight to these considerations that we do.

Suppose we believe that Potter unjustifiably injured George by a certain action, and that while Potter willingly acted with the intent of injuring George, he could not have been expected to omit this action because, through no fault of his own, he had no internal reason to do so. If we regard Potter's treatment of George as unjustified, we at least think that Potter does not agree with us about the status of things that are *for us* important normative considerations. George, in particular, may think that his welfare is valuable, and that a person of good will who is appropriately sensitive to value will see a reason to avoid actions that are contrary to his welfare.

From George's perspective, then, since Potter has failed to see his welfare as reason-giving, this suggests that he is not disposed toward George as a person of good will would be.

Because of George's judgment about the importance of his own welfare, Potter's action appears unjustifiable and offensive to George. Potter, of course, will see things differently: he thinks his treatment of George is entirely appropriate and that he has no reason to refrain from it. This may mean that Potter has no internal reason to omit his action, but regardless of whether this is the case, it is still true that Potter does not view the prospect of George's injury as a reason to refrain from his action. Thus, Potter's action expresses the offensive judgment that George's welfare is not particularly valuable.¹⁵ And even if George agrees that Potter had no (internal) reason to refrain from his action, it is inappropriate to demand that George regard Potter's action as unobjectionable because this asks George to concede that his welfare is not normatively significant. Even if George agrees that, in the internalist sense, Potter has no reason to care how he fares, it is still appropriate for George to insist that his welfare is valuable and to see Potter's rejection of this claim as a manifestation of ill will.

If we agree with George that his welfare is valuable, and that a person of good will would see his welfare as a source of reasons, we should conclude that Potter's judgment about the significance of George's welfare reasonably elicits blaming responses on George's part. By contrast, if we think that George's welfare has little normative significance, we are likely to find George's blame inappropriate. However, this conclusion would stem not from the thought that

¹⁵ What if a wild animal injured George? Would we say that the animal's behavior expresses a blame-grounding rejection of the significance of George's welfare? I think not. The behavior of animals is not morally significant in this way because a judgment like, "George's injuries don't matter," is not meaningfully attributed to them. I don't mean that nonhuman animals can never be described as being sensitive to reasons but rather that their behavior does not have the same significance for us as that of beings whose behavior is more richly and generally informed by evaluative judgments.

Potter can't reasonably be expected to recognize the significance of George's welfare, but rather from our disagreement with George about this significance.

It may be thought that Potter can reject the normative status of George's welfare in a morally relevant way only if he had rational access to an accurate judgment about this status.¹⁶ I don't see why this should be so. What reason is there to think that the expressive significance, for George and for us, of Potter's behavior should hang on whether Potter might rationally have made different judgments about reasons? After all (and as I pointed out above), Potter's actual behavior, and his actual failure to agree with George about the significance of his welfare, need not have been caused by a glitch, but may be a thoroughly deliberate and controlled exercise of his agency. And suppose that Potter *did* have some internal reason to take George's welfare as a constraint on his behavior, but that he still unjustifiably and deliberately injured him. If we did not already think that Potter's action was offensive in a blame-grounding way, I do not see how adding this element to the story would make blame appropriate. What matters is that George believes (and we believe) that he has standing that makes Potter's treatment of him illegitimate. Regardless of whether Potter could have been rationally moved to accept our view about the treatment to which George is entitled, his knowing and willing behavior demonstrates that he rejects this view.

In one of his discussions of internal reasons, Bernard Williams characterizes blame as functioning like advice given after the fact—it tells wrongdoers what they ought to have done. As Williams says, “if ‘ought to have’ is appropriate afterwards in the modality of blame, then (roughly) ‘ought to’ was appropriate at the time in the modality of advice” (Williams 40). However,

¹⁶ This would be similar to Gary Watson's claim that a wrongdoer flouts a moral demand in a way that justifies resentment only if he is capable of recognizing the validity of the demand (Watson 2004 234).

‘ought to’ in the modality of advice implies ‘can,’ because advice aims to offer something as a candidate for a deliberative conclusion. If ϕ -ing is not available to the agent, ‘You ought to ϕ ’ cannot function as a piece of advice about what he should now do; when it is a matter of what I am to do, manifestly ‘I cannot’ acts as a stopper (Williams 40).

Similarly, we might worry that if a wrongdoer has no internal reason to refrain from an action, this is a “stopper” on blame because there is no point to the advice that is supposedly implicit in blame.

I agree that when we blame a person, we are typically committed to the claim that she ought not to have done what she did. But the “ought” of blame is not always, or at least not solely, in the mode of advice. It can also be an “ought” that points to an ideal or to a moral fact: the fact, for example, that George deserves better treatment from Potter, and that a person of good will would not have disregarded George’s interests and welfare. This sort of “ought” does not imply “can.” Of course, as a form of advice, it may be futile for George to insist to Potter that his welfare matters. But as a form of *protest* this insistence is a natural way of expressing the moral offense and resentment involved in blame.¹⁷ The reasons that matter most for blame are *our* reasons—the considerations recognized by the victim and by those who sympathize with her. Whether these considerations can also be reasons for the blamed party is, I think, of secondary importance.¹⁸

¹⁷ I argue that moral blame is sometimes best construed as a form of moral protest in (2012a).

¹⁸ A referee for Oxford University Press suggests that I might note the affinity between what I say in this section and, e.g., the sort of view once defended by Philippa Foot (1972). There certainly seem to me to be affinities here but I am not sure how to characterize them; I will simply note that I am very attracted to the sort of view Foot outlines. Another referee suggests that I can in fact say that it is reasonable to expect wrongdoers like Mr. Potter to correct their moral ignorance. This suggestion relies on my characterization of these wrongdoers as expressing ill will, the fact that expressions of ill will (arguably) flout the demand to treat others with reasonable regard, and the fact that this demand is connected to our normative expectations of other agents. I don’t doubt that this proposal is workable but (in this and previous work) I prefer to grant to opponents that certain moral demands and expectations are off the table in Potter’s case (and other related cases). I think this helps to make the debate more clear than it would be otherwise. However, the referee’s comment prompts me to note that my disagreement with FitzPatrick is in some ways less stark than I presented it in Section 2. FitzPatrick and I agree that Potter can reasonably be held to our moral standards in the sense that we regard him as a proper target of moral blame. (I suspect this is the sense in which the referee believes I can say that we properly expect better things of Potter.) However, for FitzPatrick, the

6. Circumstantial Knowledge And Blame

On the account I have presented, when it comes to thinking about how knowledge relates to blameworthiness, we should consider what a wrongdoer needs to know in order for her actions to express the attitudes and judgments that make blame appropriate. What does an agent need to know before we attribute to her a judgment like, “your objections and interests can be overlooked when I am deciding how to act”? As I have argued, awareness that one’s actions are wrong is not necessary for the presence or expression of such judgments. Much more relevant in this context is knowledge of the effects that one’s actions will have on others. If a person knows that her action will injure someone who objects to this result, then it is reasonable to attribute to her the judgment that the other’s injuries and objections are not a decisive reason to refrain from the action. Judgments like this help to ground blaming responses because of the way they call into question the standing of the other to raise objections to certain forms of treatment and to cite her interests as normative considerations.

Of course, we will not always think that an agent is open to blame when her actions express the judgment that another’s interests and objections may be overlooked. As I suggested with my example of Robert E. Lee in Section 4, our views about an agent’s blameworthiness vary with the degree to which we accept the judgments that move her. Most people believe that it is permissible to use violence in self-defense against an aggressor. If we take an instance of self-defense to be justified, then we likely think that whatever objections the aggressor might raise to the force used against him are rightly overlooked. Thus, we would be unlikely to blame the person who defended herself even if she acted on the judgment that, in the relevant instance, the aggressor’s welfare can be overlooked. The injured aggressor may resent the one who injured

legitimacy of blaming Potter depends on his having been able to avoid his moral ignorance. On my view, Potter’s blameworthiness does not depend on this but only on the actual judgments that inform his behavior.

him in self-defense, but by our lights, and by those of the one who acted in self-defense, this blame will appear inapt and unreasonable.

Similarly, if I uproot a plant, my action may express an implicit judgment that plants do not have standing to raise objections to the way I treat them. But I am not blameworthy here because this judgment is accurate; treating plants this way *is* unobjectionable. Whatever objections there may be to uprooting a plant are not grounded in the fact that plants are, in themselves, the sort of things that can be treated objectionably.

Suppose, however, that I am wrong about this, and that plants can suffer, and that they have a perspective from which to raise objections to my treatment of them. In the context of criticizing views like the one defended in this paper, Neil Levy proposes just this case:

Suppose that there is a kind of harm that is objectively morally relevant, but of which we are ignorant. Suppose, for instance, that plants can be harmed, and that this harm is a moral reason against killing or treading on them. In that case, many of us are causally responsible for a great many moral harms. Are we morally responsible for them? Do we flout a moral requirement, and challenge plants' standing as objects to which some moral consideration is owed? No to all these questions: If we do not grasp the moral requirement, and this ignorance is not culpable, we do nothing blameworthy (Levy 2005 9).¹⁹

Now I agree that we are not blameworthy for stepping on plants in Levy's example, but the example is offered as an instance of exculpation by way of nonculpable *moral* ignorance and this is not what it shows. Levy's example fails to illustrate the exculpatory power of moral ignorance because what explains why we are not blameworthy in the example is that we lack crucial information about how walking on plants affects them. Since I do not know that plants can be harmed, my stepping on one does not express a denial of the significance of its being

¹⁹ I should note that Levy is considering psychopaths (rather than relatively normal wrongdoers like Mr. Potter) when he offers this example, but I don't think this affects its usefulness here. David Shoemaker (2011) treats a similar example; I reply to Shoemaker in (2012b).

harm. So on the view I advocate, stepping on a plant does not express a judgment that could properly ground blame.

The problem is that Levy's example does not involve the relevant sort of moral ignorance. There is, of course, a kind of moral ignorance in the example: we lack knowledge of the moral status of stepping on plants (because we are ignorant of their capacity for suffering). But this moral ignorance derives from circumstantial ignorance and not from ignorance of a moral principle. In the example, I presumably know that I generally have moral reasons to not cause pain; I just don't know that stepping on plants causes them pain. However, the moral ignorance of someone like Mr. Potter is just the reverse: he knows that his actions will cause George Bailey and others pain, but he does not know (because he does not believe) that this counts decisively against performing these actions.

For Levy's example to help us see that Potter's sort of moral ignorance counts as an excuse, the example would have to feature Potter's sort of ignorance. The example would need to be one in which I know, say, that stepping on a plant will cause it pain and I (nonculpably, but wrongly) believe that this pain does not matter. Now if we think that I am blameless in this revised version of Levy's example, we might draw a conclusion about the exculpatory significance of nonculpable ignorance of moral principles. Of course, I would not agree that I am blameless in this revised case. If I knowingly (and unjustifiably) cause a plant severe pain, then I am open to resentment and indignation because my action dismisses the normative significance of the plant's pain, and this is something to which the plant (or at least one who is concerned for its welfare) could reasonably object.

Conclusion

I have argued against the assumption that morally ignorant wrongdoers are open to moral blame only if they are culpable for their ignorance. Thus, I reject skepticism about moral responsibility that depends on this assumption. On the view I have defended, the attitudes involved in moral blame are responses to the features of an action that make it objectionable from the perspective of the one who issues the blame. One important way that an action can appear objectionable to us is that it expresses a judgment with which we disagree about the significance of the needs and interests of those affected by the action. Whether a wrongdoer's action has this feature depends more on whether she is aware of the consequences of her behavior than on whether she regards her behavior as wrong.

References

- Foot, Philippa. 1972. "Morality as a System of Hypothetical Imperatives." *The Philosophical Review* 81: 305-316.
- FitzPatrick, William J. 2008. "Moral Responsibility and Normative Ignorance: Answering a New Skeptical Challenge." *Ethics* 118: 589-613.
- Frankfurt, Harry. 1969. "Alternate Possibilities and Moral Responsibility." *Journal of Philosophy* 66: 829-39.
- Levy, Neil. 2005. "The Good, the Bad, and the Blameworthy." *Journal of Ethics and Social Philosophy* 1: 1-16.
- 2009. "Culpable Ignorance and Moral Responsibility: A Reply to FitzPatrick." *Ethics* 119: 729-741.
- 2011. *Hard Luck: How Luck Undermines Free Will and Moral Responsibility*. New York: Oxford University Press.
- Rosen, Gideon. 2003. "Culpability and Ignorance." *Proceedings of the Aristotelian Society* 103: 61-84.
- 2004. "Skepticism about Moral Responsibility." *Philosophical Perspectives* 18: 295-313.

- 2008. "Kleinbart the Oblivious and Other Tales of Ignorance and Responsibility." *The Journal of Philosophy* 105: 591-610.
- Scanlon, T. M. 1998. *What We Owe to Each Other*. Cambridge, MA: Harvard University Press.
- 2008. *Moral Dimensions: Permissibility, Meaning, Blame*. Cambridge, MA: Harvard University Press.
- Shoemaker, David. 2011. "Attributability, Answerability, and Accountability: Toward a Wider Theory of Moral Responsibility." *Ethics* 121: 602-632.
- Smith, Angela. 2005. "Responsibility for Attitudes: Activity and Passivity in Mental Life," *Ethics* 115: 236-271.
- 2008. "Control, Responsibility, and Moral Assessment." *Philosophical Studies* 138: 367-392.
- Smith, Holly. 1983. "Culpable Ignorance." *The Philosophical Review* 92: 543-571.
- Talbert, Matthew. 2008. "Blame and Responsiveness to Moral Reasons: Are Psychopaths Blameworthy?" *Pacific Philosophical Quarterly* 89: 516-535.
- 2012a. "Moral Competence, Moral Blame, and Protest." *The Journal of Ethics* 16: 89-109.
- 2012b. "Aliens, Accountability, and Psychopaths: A Reply to Shoemaker." *Ethics* 122: 562-574.
- Watson, Gary. 2004. "Responsibility and the Limits of Evil: Variations on a Strawsonian Theme." In *Agency and Answerability*, 219-259. New York: Oxford University Press.
- Weingarten, Gene. 2008. "Fatal Distraction." *The Washington Post*, March 8, W8.
- Williams, Bernard. 1995. "Internal Reasons and the Obscurity of Blame." In *Making Sense of Humanity*, 35-45. Cambridge: Cambridge University Press. 1995.
- Zimmerman, Michael. 1997. "Moral Responsibility and Ignorance." *Ethics* 107: 410-426.
- 2008. *Living with Uncertainty*. Cambridge: Cambridge University Press.